re:Invent

Unmeltable Infrastructure at Scale: Using Apache Kafka, Twitter Storm, and ElasticSearch on AWS

Jim Nisbet
CTO and VP of Engineering, Loggly

Philip O'Toole

Lead Developer, Infrastructure, Loggly

November 2013



What Loggly Does

- Log Management as a service
 - Near real-time indexing of events
- Distributed architecture, built on AWS
- Initial production services in 2010
 - Loggly Generation 2 released in Sept 2013
- Thousands of customers





Agenda for this Presentation

- A bit about logging
- Lessons learned from our first generation
- How we leverage AWS services
- Our use of Kafka, Storm, ElasticSearch
- What worked well for us and what did not

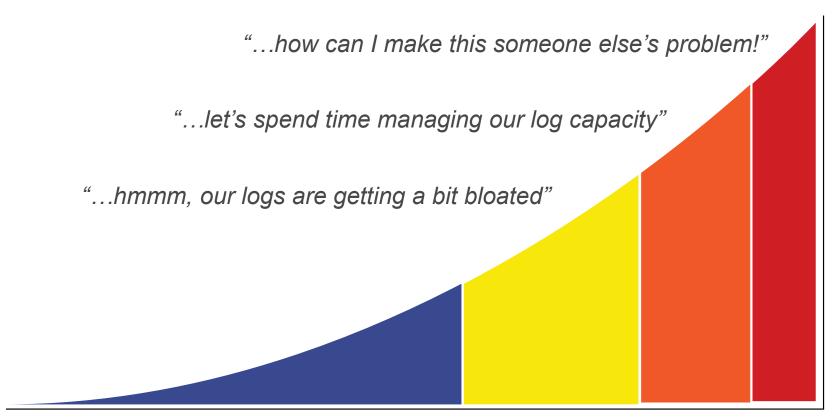


Log Management

- Everyone starts with
 - A bunch of log files (syslog, application specific)
 - On a bunch of machines

- Management consists of doing the simple stuff
 - Rotate files, compress and delete
 - Information is there but awkward to find specific events
 - Weird log retention policies evolve over time





Log Volume



Best Practices in Log Management

- Use existing logging infrastructure
 - Real time syslog forwarding is built in
 - Application log file watching
- Store logs externally
 - Accessible when there is a system failure
- Log messages in machine parsable format
 - JSON encoding when logging structured information
 - Key-value pairs



From the Trenches...

- Managing Applications vs. Managing Logs
 - Do not make this is an either/or proposition!



If you get a disk space alert, first login...

% sudo rm -rf /var/log/apache2/*

Admit it, we've all seen this kind of thing!



You Have Logs...

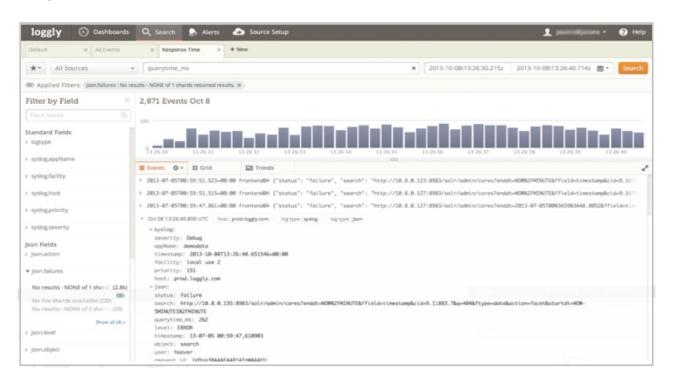
2013-10-25T18:35:43.387+0000: 441.482: [GC [PSYoungGen: 2430541K->268617K(2484544K)] 7687523K->5660738K(8076992K), 0.3266870 secs] [Times: user=1.05 sys=0.17, real=0.33 secs] 2013-10-25T18:35:43.714+0000: 441.809: [Full GC [PSYoungGen: 268617K->0K(2484544K)] [ParOldGen: 5392121K->354965K(5592448K)] 5660738K->354965K(8076992K) [PSPermGen: 44444K->44395K(83968K)], 0.9225290 secs] [Times: user=2.22 sys=0.26, real=0.92 secs]

 In this case, JVM garbage collection logs enabled with...

-XX:+PrintGCDetails -XX:+PrintGCDateStamps -XX:+PrintGCTimeStamps

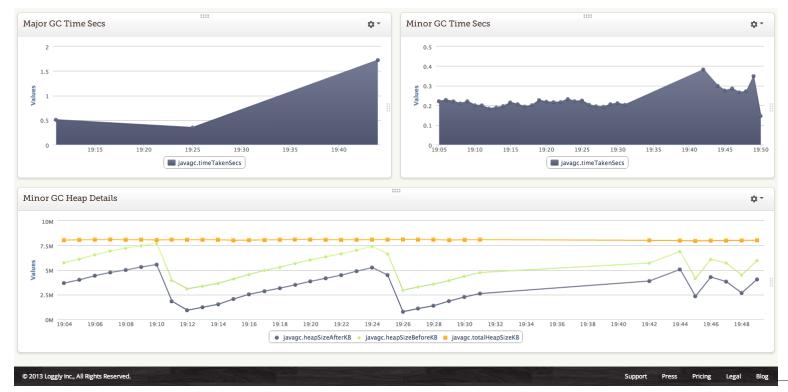


Yes, you need to search these logs





But you also need to to spot trends





Loggly Offers Logging as a Service



Loggly First Generation

- Logging as a service
 - Near real-time searchable logs
- Thousands of customers
 - Transmission rates from 10 events/sec to 100k events/sec
 - When customers systems are busy they send more logs
 - Log traffic has distinct bursts; bursts can last for several hours
- Amazon EC2 deployment
 - We used EC2 Instance storage
- SOLR Cloud
 - Full power of Lucene search
 - Tens of thousands of shards (with special 'sleep shard' logic)
- ZeroMQ for message queue



First Generation Lessons Learned

- Event ingestion too tightly coupled to indexing
 - Manual re-indexing for temporary SOLR issues
- Multiple Indexing strategies needed
 - 4 orders of magnitude difference between our high volume users and our low volume users (10 eps vs. 100,000+ eps)
 - Too much system overhead for low volume users
 - Difficult to support changing indexing strategies for a customer



Big Data Infrastructure Solutions

We are <u>not</u> alone...

- Our challenges
 - Massive incoming event stream
 - Fundamentally multi-tenant
 - Scalable framework for analysis
 - Near real-time indexing
 - Time series index management





Apache Kafka

Overview

- An Apache project initially developed at LinkedIn
- Distributed publish-subscribe messaging system
- Specifically designed for real time activity streams
- Does not follow JMS Standards nor uses JMS APIs

Key Features

- Persistent messaging
- High throughput, low overhead
- Uses ZooKeeper for forming a cluster of nodes
- Supports both queue and topic semantics





Message Queue Performance

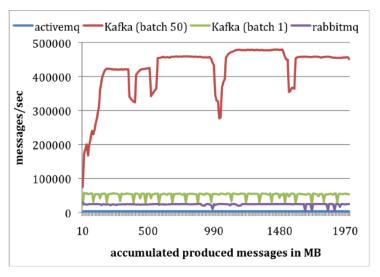


Figure 4. Producer Performance

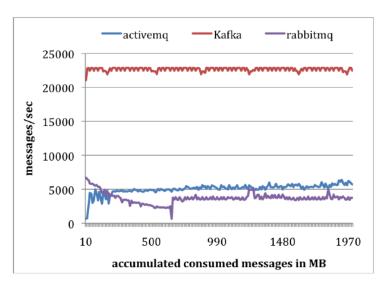


Figure 5. Consumer Performance

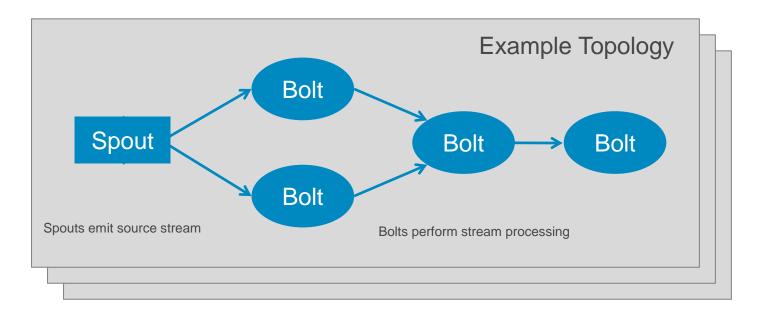
http://research.microsoft.com/en-us/um/people/srikanth/netdb11/netdb11papers/netdb11-final12.pdf

Storm Framework

- Storm (open sourced by Twitter)
 - Open sourced September 2011
 - Now an Apache Software Foundation project
 - Currently Incubator Status
- Framework is for stream processing
 - Distributed
 - Fault tolerant
 - Computation
 - Fail-fast components



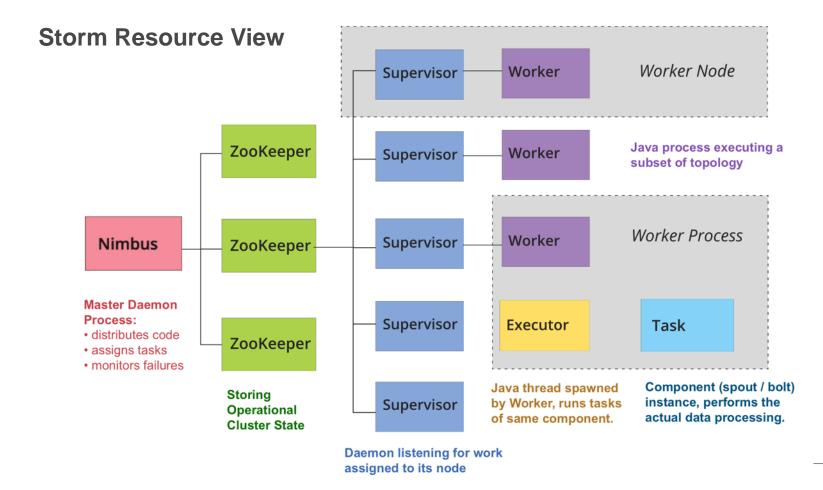
Storm Logical View



Storm terminology

• Streams, Spouts, Bolts, Tasks, Workers, Stream Groups and Topologies







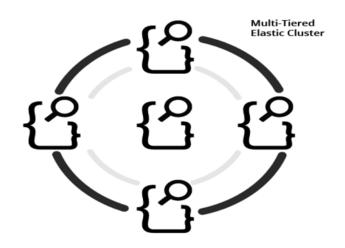
ElasticSearch

- Open source
 - Commercial support available from ElasticSearch.com
 - Growing open-source community
- Distributed search engine
- Fully exposes Lucene search functionality
- Built for clustering from the ground-up
- High availability
- Multi-tenancy



ElasticSearch In Action

- Add/delete nodes dynamically
- Add indices with REST API
- Indices and Nodes have attributes
 - Indices automatically moved to best Nodes
- Indices can be sharded
- Supports bulk insertion of events
- Plugins for monitoring cluster





Our Second Generation





Generation 2 – The Challenge

- Always accept log data
 - Never make a customer's incident worse
- Never drop log data
 - A single log message could be critical
- True Elasticity



Perfect Match For Real Time Log Events

- Apache Kafka
 - Extremely high-performance pub-sub persistent queue
- Consumer tracks their location in queue
 - A good fit for our use cases
- Multiple Kafka brokers
 - Good match for AWS
 - Multiple brokers per region
 - Availability Zone separation

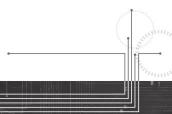


Real Time Event Processing

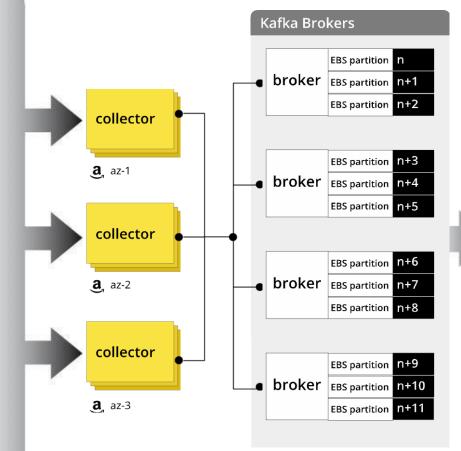
- Twitter Storm
 - Scalable real-time computation system
- Storm used as a "pull" system
 - Provisioned for average load, not peak load
 - Input from Kafka queue
 - Worker nodes can be scaled dynamically
- Elasticity is key
 - Another good match for AWS
 - Able to scale workers up and down dynamically



Log Event Ingestion









event ingestion

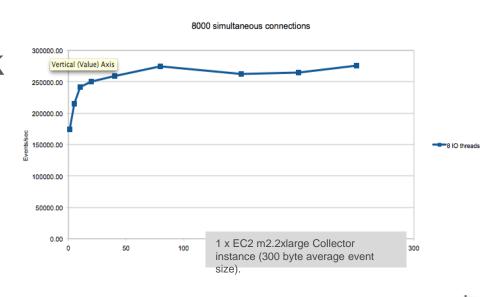
+100ks / second

syslog protocol (TCP/UDP)

HTTP

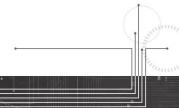
Loggly Collector Performance

- C++ multi-threaded
- Boost ASIO framework
- Each Collector can handle 250k+ events per second
 - Per m2.2xlarge instance



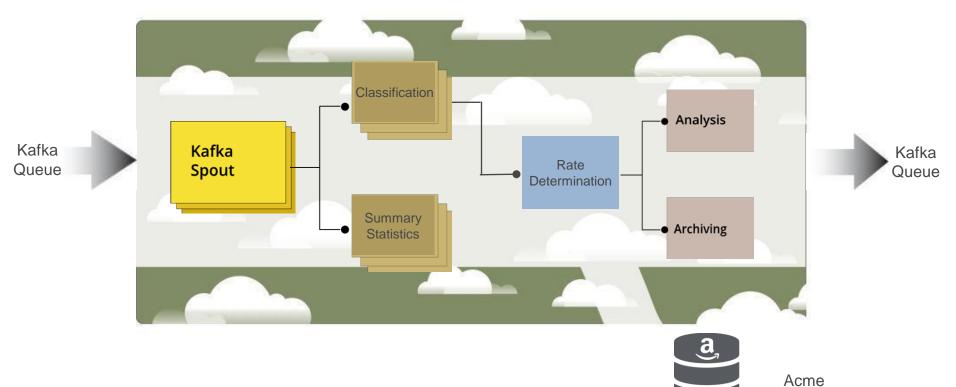


Processing Events





Storm Event Processing



S3 Bucket

Event Pipeline in Summary

- Storm provides Complex Event Processing
 - Where we run much of our secret-sauce
- Kafka contains both raw and processed Events
- Snapshot the last day of Kafka events to S3





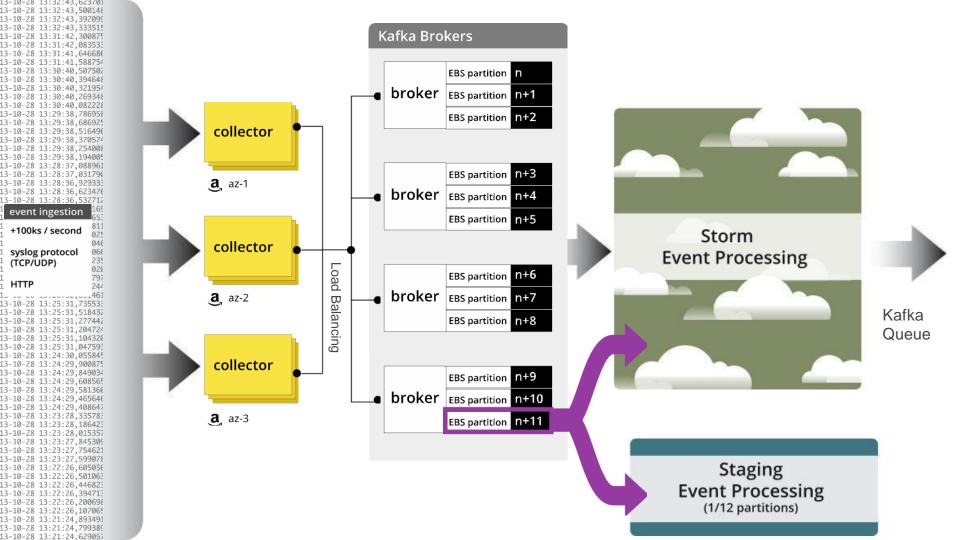
Loggly and Index Management

- Indices are time-series data
 - Separated by customer
 - Represent slices of time
 - Higher volume index will have shorter time slice
- Multi-tier architecture for efficient indexing
 - Multiple indexing tiers mapped to different AWS instance types
- Efficient use of AWS resources



Staging Pre-Production System



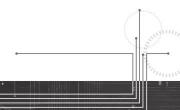


Kafka enables Staging Architecture

- Kafka Broker doesn't care if there are multiple consumers
- Staging system runs pre-production code
- Pub-sub allows us to randomly index a fraction of our production load
- A highly-effective pre-production system



AWS Deployment Details





AWS Deployment Instances – Collection

Collector

c1.xlarge

- Compute-optimized
- High-traffic ingestion points
- Disk not important



m2.2xlarge

- Memory-optimized
- Disk buffer caching



4K Provisioned IOPs EBS

- Ensures consistent IO
- No noisy-neighbors
- Persistent storage



AWS Deployment Instances – Processing

Storm Supervisor

c1.xlarge

- Compute-optimized
- CPU-intensive processing
- Network IO

Storm Nimbus

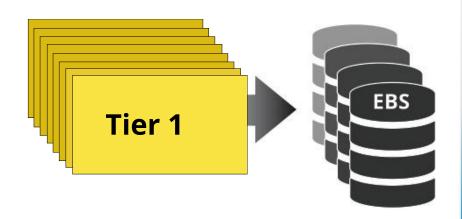
ZooKeeper

m1.xlarge

- General-purpose
- Configuration
- Management

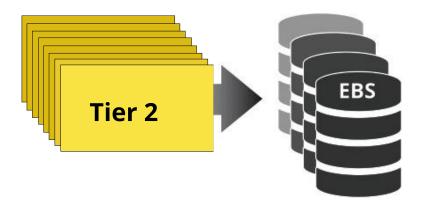


AWS Deployment Instances – Indexing



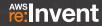
cc2.8xlarge

4K Provisioned IOPs EBS

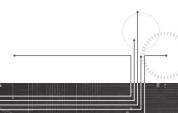


m2.4xlarge

4K Provisioned IOPs EBS



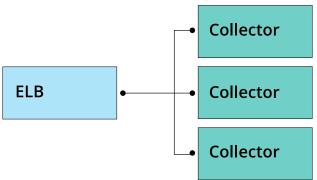
A Few False Starts





Elastic Load Balancing in front of Collector Had Limitations

Initial testing used Elastic Load Balancing for incoming events:



- Elastic Load Balancing doesn't allow forwarding port 514 (syslog)
- Elastic Load Balancing doesn't support forwarding UDP
- Event traffic can burst and hit Elastic Load Balancing performance limits



Amazon Route 53 DNS Round Robin a Win

- DNS Round Robin is pretty basic load balancing
 - Not a bump in the wire
- Take advantage of AWS failover health checks
 - When a collector goes out of service, it will be out of the DNS rotation
- Round Robin across multiple regions, AZs
 - Latency based resolution optimizes inbound traffic





Our First Plan for Log Events

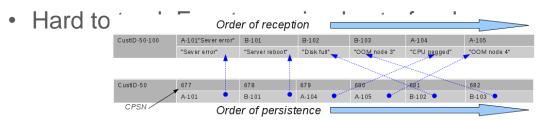
- Cassandra
 - Highly scalable key-value store
 - Impressive write performance a good match for us
 - Apache project plus commercial support with DataStax

- Use Cassandra for both our Event Queue and Persistent Store
 - Our strategy was to get the raw event in to Cassandra
 - ...then perform workflow processing on events



Design meets Reality

Cassandra not designed to be a message queue



- Multi-tenancy requires handling data bursts
 - Collectors still needed to be able to buffer to disk
 - Added complexity and became a point of failure

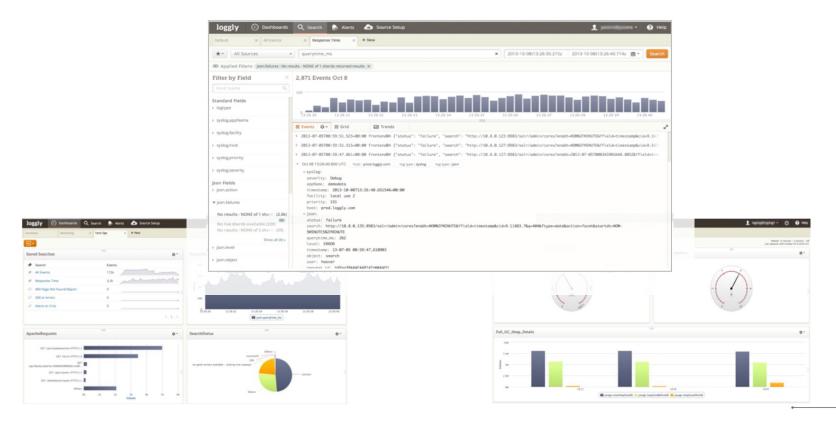


Big Wins

- Leveraging AWS services
 - Multi-region, multi-AZ
 - Provisioned IOPS for availability and scale
 - Amazon Route 53 DNS support with latency resolution
 - Easy to increase and decrease Storm resources
- Leveraging Open Source infrastructure
 - Apache Kafka
 - Twitter Storm
 - ElasticSearch
- Pre-production "Staging" system



The Means to an End





Feedback

• Questions?

Jim Nisbet (niz@loggly.com)
CTO and VP of Engineering, Loggly

Philip O'Toole (philip@loggly.com)
Lead Developer, Infrastructure, Loggly



Follow us @loggly!



re:Invent

Please give us your feedback on this presentation

ARC303

As a thank you, we will select prize winners daily for completed surveys!



